

ford-car-price-prediction

July 19, 2023

Project: **Ford car price prediction**

Algorithm Used: **Multiple regression, Random forest regression, Gradient boost regressor and XGboost regressor**

Steps:

- Loading the dataset using pandas library
- preprocessing the dataset.
- Exploratory Data Analysis
- Model building using Multiple regression Algorithm, Random forest regressor, Gradient boost regressor and XGboost regressor
- Model testing using MSE, MAE, R2-SCORE, RMSE

```
[ ]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import plotly.express as px
import warnings
warnings.filterwarnings('ignore')
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.ensemble import AdaBoostRegressor
from sklearn.ensemble import GradientBoostingRegressor
from xgboost import XGBRegressor
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
df=pd.read_csv('/content/ford.csv')
df
```

```
[ ]:      model  year  price  transmission  mileage  fuelType  tax  mpg  \
0      Kuga   2019  18990      Manual      8389   Petrol   150  35.3
1    Fiesta   2019  21999      Manual      4000   Petrol   145  40.3
2      Ka+    2020  11999      Manual      2000   Petrol   145  43.5
```

3	KA	2018	9899	Manual	6000	Petrol	145	43.5
4	Ka+	2018	9999	Manual	15000	Petrol	145	43.5
...
17961	Mustang	2017	27890	Semi-Auto	26452	Petrol	580	23.5
17962	Mustang	2020	42999	Manual	10	Petrol	145	23.7
17963	Mustang	2020	48000	Manual	50	Petrol	145	23.9
17964	Mustang	2020	40495	Semi-Auto	3200	Petrol	145	24.8
17965	Mondeo	2017	15499	Automatic	10162	Petrol	235	38.2

```

engineSize
0          0.0
1          0.0
2          0.0
3          0.0
4          0.0
...
17961      5.0
17962      5.0
17963      5.0
17964      5.0
17965      5.0

```

[17966 rows x 9 columns]

Data Preprocessing

```
[ ]: df.dtypes
```

```

[ ]: model          object
     year          int64
     price          int64
     transmission   object
     mileage        int64
     fuelType       object
     tax            int64
     mpg            float64
     engineSize     float64
     dtype: object

```

```
[ ]: df.isna().sum()
```

```

[ ]: model          0
     year           0
     price          0
     transmission   0
     mileage        0
     fuelType       0

```

```
tax          0
mpg          0
engineSize   0
dtype: int64
```

```
[ ]: df.duplicated().sum()
```

```
[ ]: 154
```

```
[ ]: #dropping duplicate values
df1 = df.drop_duplicates().reset_index(drop=True)
```

```
[ ]: df1.duplicated().sum()
```

```
[ ]: 0
```

```
[ ]: df1.shape
```


```
[ ]: (17812, 9)
```

```
[ ]: #removing negative values
df1=df1[df1['price'] >= 0]
```

Explortory Data Analysis

```
[ ]: fig, ax = plt.subplots(figsize=(6,6))
sns.distplot(df1['price'],kde=False)
```

```
[ ]: <Axes: xlabel='price'>
```




```
[ ]: a=df1.groupby('model')['price'].sum().reset_index()
      index_to_drop=23
      a=a.drop(index_to_drop)
      a1=a.sort_values(by='price',ascending=False)[['model','price']].head(10)
      a1
```

```
[ ]:
      model    price
5    Fiesta  66332283
6    Focus   60030224
13   Kuga    34928705
2    EcoSport 14052864
14   Mondeo   6276581
1     C-MAX   5373119
18   S-MAX   5190688
3     Edge   4655584
12    Ka+    4549602
```

8 Galaxy 4051648

```
[ ]: fig,ax=plt.subplots(figsize=(8,8))
sns.barplot(x='model',y='price',data=a1,ax=ax,palette='pastel')
```

```
[ ]: <Axes: xlabel='model', ylabel='price'>
```

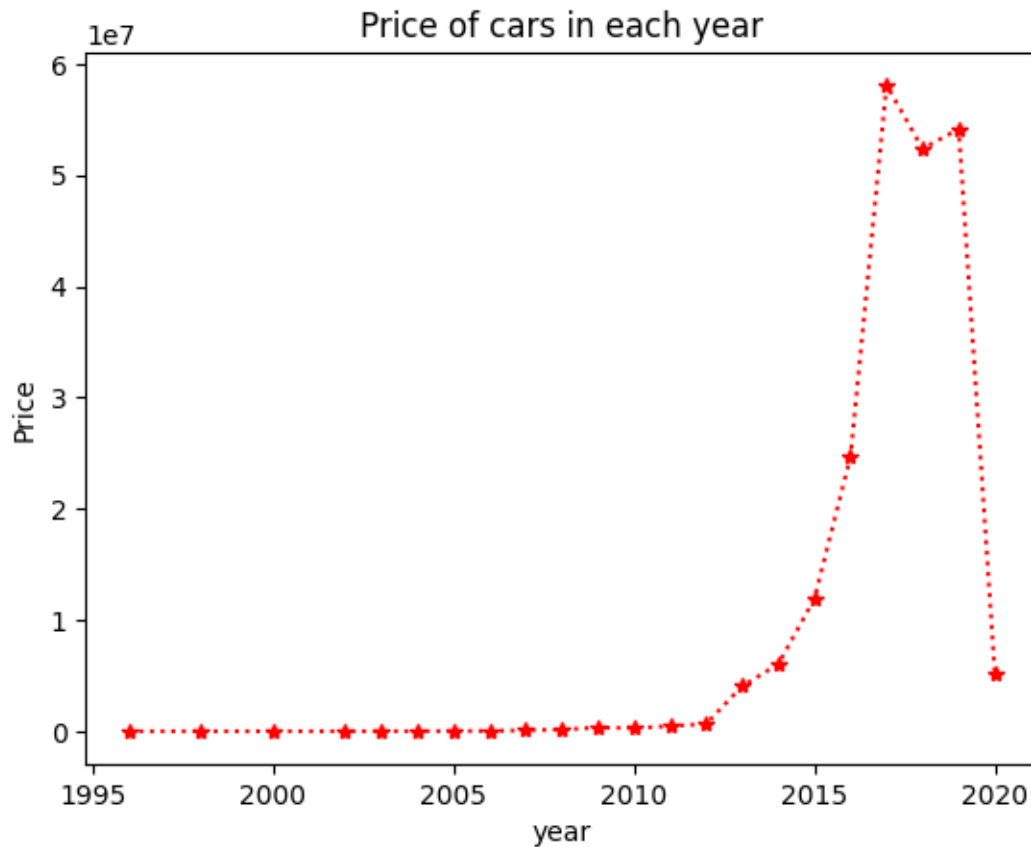


```
[43]: b=df1.groupby('year')['price'].sum().reset_index()
b1=b.sort_values(by='year')
b1=b1.drop(22)
b1
```

```
[43]:
```

	year	price
0	1996	3000
1	1998	2699
2	2000	1995
3	2002	5785
4	2003	6189
5	2004	5744
6	2005	25488
7	2006	28634
8	2007	83314
9	2008	146342
10	2009	338495
11	2010	271903
12	2011	472113
13	2012	653307
14	2013	4061935
15	2014	6021380
16	2015	11908049
17	2016	24677121
18	2017	58082890
19	2018	52355498
20	2019	54139409
21	2020	5247552


```
[45]: x=b1['year']  
y=b1['price']  
plt.plot(x,y,"*:r")  
plt.xlabel("year")  
plt.ylabel("Price")  
plt.title("Price of cars in each year")  
plt.show()
```



```
[ ]: c=df1.groupby('transmission')['price'].sum().reset_index()
      c1=c.sort_values(by='price')
      c1
```

```
[ ]:   transmission    price
      2   Semi-Auto  16015323
      0   Automatic  21262532
      1    Manual   181267482
```

```
[ ]: y=c1['price']
      labels=['Semi-Auto','Automatic','Manual']
      myexplode=[0,0,0.2]
      plt.pie(y,labels=labels,explode = myexplode,shadow=True,autopct='%1.1f%%')
      plt.show()
```




```
[ ]: d=df1.groupby('fuelType')['tax'].sum().reset_index()
      d1=d.sort_values(by='tax')
      d1
```

```
[ ]:   fuelType    tax
      1  Electric      0
      3    Other      0
      2   Hybrid   2215
      0   Diesel  580305
      4   Petrol 1435847
```

```
[47]: size=[200, 400, 600, 800, 1000]
      color=[1, 2, 3, 4, 5]
      px.scatter(d1, x='fuelType', y='tax',size=size, color=color)
```

```
[ ]: sns.jointplot(data=df1, x=df1['mileage'], y=df1['price'],color='g')
```

```
[ ]: <seaborn.axisgrid.JointGrid at 0x7d735ab0b550>
```

```
[ ]: sns.heatmap(df1.corr(),annot=True,linewidths=1)
```

```
[ ]: <Axes: >
```



Observations

- More number of cars were purchased in the price range **10000** Dollars to **20000** Dollars
- **Fiesta**, **Focus** and **kugo** are the cars in that comes under highest price range in the Ford cars list.
- The prices were high between the years **2015** to **2018**.
- **Manual** cars are highly priced than automatic and semi auto
- Cars that run on **petrol** gives highest tax followed by diesel and **Electric** does cars doesnot have to pay tax
- As the mileage increases price decreases.

Model Building

```
[ ]: X=df1.drop(['price'],axis=1)
     y=df1['price']
```

```
[ ]: le=LabelEncoder()
     lst=['model','transmission','fuelType']
     for i in lst:
         X[i]=le.fit_transform(X[i])
```

```
[ ]: X
```

```
[ ]:      model  year  transmission  mileage  fuelType  tax  mpg  engineSize
0         13  2019             1     8389         4  150  35.3         0.0
1          5  2019             1     4000         4  145  40.3         0.0
2         12  2020             1     2000         4  145  43.5         0.0
3         11  2018             1     6000         4  145  43.5         0.0
4         12  2018             1    15000         4  145  43.5         0.0
...      ...  ...             ...      ...      ...      ...      ...
17807      15  2017             2    26452         4  580  23.5         5.0
17808      15  2020             1         10         4  145  23.7         5.0
17809      15  2020             1         50         4  145  23.9         5.0
17810      15  2020             2     3200         4  145  24.8         5.0
17811      14  2017             0    10162         4  235  38.2         5.0
```

[17812 rows x 8 columns]

```
[ ]: #scaling using standard scaler
ms=MinMaxScaler()
X_ms=ms.fit_transform(X)
```

```
[ ]: #Performing train_test_split
X_train,X_test,y_train,y_test=train_test_split(X_ms,y,test_size=0.
↪3,random_state=0)
```

```
[ ]: #model building using Multiple regression
mlr=LinearRegression()
mlr.fit(X_train,y_train)
y_pred=mlr.predict(X_test)
y_pred
```

```
[ ]: array([12548.90861712,  9106.24387537,  8051.36725607, ...,
          11849.31347771,  8568.45511522, 12324.32982938])
```

Multiple Regression

```
[ ]: #model validation
import numpy as np
print("mean absolute error:",mean_absolute_error(y_test,y_pred))
print("mean squared error:",mean_squared_error(y_test,y_pred))
print("root mean squared error:",np.sqrt(mean_squared_error(y_test,y_pred)))
print("r2-score:",r2_score(y_test,y_pred))
```

```
mean absolute error: 1758.779360016115
mean squared error: 6758138.433333946
root mean squared error: 2599.64198176094
r2-score: 0.7083660598390757
```

Random Forest Regressor

```
[ ]: rs=RandomForestRegressor()  
rs.fit(X_train,y_train)  
y_pred1=rs.predict(X_test)  
y_pred1
```

```
[ ]: array([11854.45, 8205.83, 8469.2 , ..., 12626.35, 7742.25, 12058.46])
```

```
[ ]: print("mean absolute error:",mean_absolute_error(y_test,y_pred1))  
print("mean squared error:",mean_squared_error(y_test,y_pred1))  
print("root mean squared error:",np.sqrt(mean_squared_error(y_test,y_pred1)))  
print("R2 score:",r2_score(y_test,y_pred1))
```

```
mean absolute error: 889.8758965112929  
mean squared error: 1778514.9257368413  
root mean squared error: 1333.6097351687417  
R2 score: 0.9232517474236209
```

Gradient Boosting Regressor

```
[ ]: gb=GradientBoostingRegressor()  
gb.fit(X_train,y_train)  
y_pred3=gb.predict(X_test)  
y_pred3
```

```
[ ]: array([12024.84664223, 8764.37171307, 8592.75237629, ...,  
12368.23635446, 7234.98658503, 13614.79528895])
```

```
[ ]: print("mean absolute error:",mean_absolute_error(y_test,y_pred3))  
print("mean squared error:",mean_squared_error(y_test,y_pred3))  
print("root mean squared error:",np.sqrt(mean_squared_error(y_test,y_pred3)))  
print("R2 score:",r2_score(y_test,y_pred3))
```

```
mean absolute error: 977.5392003798046  
mean squared error: 2071177.3886393986  
root mean squared error: 1439.1585696647185  
R2 score: 0.91062248449339
```

XGB Regressor

```
[ ]: xgb=XGBRegressor()  
xgb.fit(X_train,y_train)  
y_pred4=xgb.predict(X_test)  
y_pred4
```

```
[ ]: array([11974.356 , 8401.951 , 8561.761 , ..., 12414.846 , 7270.1777,  
13043.783 ], dtype=float32)
```



```
[50]: y_new=xgb.predict(ms.transform([[4,2015,1,10,4,150,40,2]]))  
      print(y_new)
```

```
[18073.955]
```